

# Memory Hierarchy

CS/COE 1541 (Fall 2020)  
Wonsun Ahn

# Experiment on kernighan.cs.pitt.edu

- The source code for the experiments are available at:  
[https://github.com/wonsunahn/CS1541\\_Fall2020/tree/master/resources/cache\\_experiments](https://github.com/wonsunahn/CS1541_Fall2020/tree/master/resources/cache_experiments)
- Or on the following directory at linux.cs.pitt.edu:  
/afs/cs.pitt.edu/courses/1541/cache\_experiments/
- You can run the experiments by doing 'make' at the root
  - It will take a few minutes to run all the experiments
  - In the end, you get two plots: IPC.pdf and MemStalls.pdf

# Four benchmarks

- `linked-list.c`
  - Traverses a linked list from beginning to end over and over again
  - Each node has 120 bytes of data
- `array.c`
  - Traverses an array from beginning to end over and over again
  - Each element has 120 bytes of data
- `linked-list_nodata.c`
  - Same as linked-list but nodes have no data inside them
- `array_nodata.c`
  - Same as array but elements have no data inside them

# Code for linked-list.c

```
#define ACCESSES 1000000000
```

```
// Define a linked list node type with no data
```

```
typedef struct node {  
    struct node* next; // 8 bytes  
    int data[30];      // 120 bytes  
} node_t;
```

```
...
```

```
// Traverse list over and over again until we've visited `ACCESSES` nodes
```

```
node_t* current = head;  
for(int i=0; i<ACCESSES; i++) {  
    if(current == NULL) current = head; // reached the end  
    else current = current->next;       // next node  
}
```

# Code for array.c

```
#define ACCESSES 1000000000
```

```
// Define a linked list node type with no data
```

```
typedef struct node {  
    struct node* next; // 8 bytes  
    int data[30];      // 120 bytes  
} node_t;
```

```
...
```

```
// Traverse array over and over again until we've visited `ACCESSES` elements
```

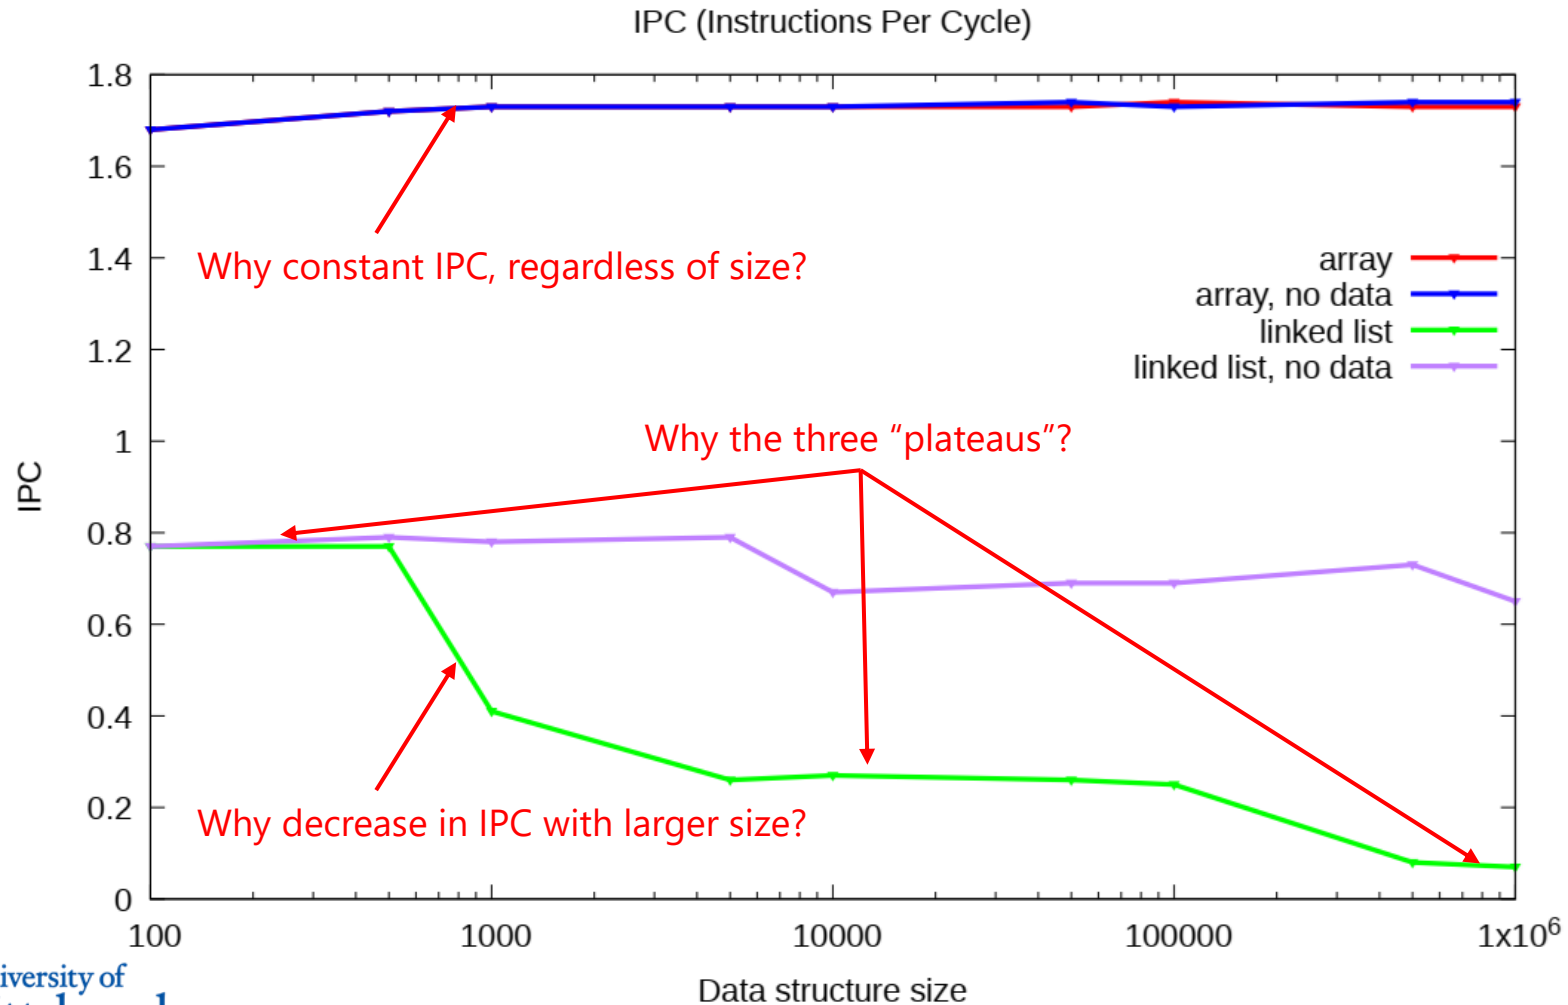
```
node_t* current = array;  
for(int i=0; i<ACCESSES; i++) {  
    if(current == array + items) current = array; // reached the end  
    else ++current; // next element  
}
```

- Two CPU sockets. Each CPU:
  - Intel(R) Xeon(R) CPU E5-2640 v4
  - 10 cores, with 2 threads per each core (SMT)
  - L1 i-cache: 32 KB 8-way set associative (per core)
  - L1 d-cache: 32 KB 8-way set associative (per core)
  - L2 cache: 256 KB 8-way set associative (per core)
  - L3 cache: 25 MB 20-way set associative (shared)
- Memory
  - 128 GB DRAM
- Information obtained from
  - "cat /proc/cpuinfo" on Linux server
  - "cat /proc/meminfo" on Linux server
  - [https://en.wikichip.org/wiki/intel/xeon\\_e5/e5-2640\\_v4](https://en.wikichip.org/wiki/intel/xeon_e5/e5-2640_v4)

# Experimental data collection


- Collected using CPU **Performance Monitoring Unit (PMU)**
  - PMU provides performance counters for a lot of things
  - Cycles, instructions, various types of stalls, branch mispredictions, cache misses, bandwidth usage, ...
- Linux **perf** utility summarizes this info in easy to read format
  - <https://perf.wiki.kernel.org/index.php/Tutorial>

# IPC (Instructions Per Cycle) Results





# Memory Stall Cycle Percentage





# Memory Technologies

---


# Static RAM (SRAM)

- SRAM uses a loop of NOT gates to store a single bit
- This is usually called a 6T SRAM cell since it uses... 6 Transistors!
- **Pros:**
  - **Very fast** to read/write
- **Cons:**
  - Volatile (loses data without power)
  - Relatively many transistors needed
    - > **expensive**



# Dynamic RAM (DRAM)

- DRAM uses **one** transistor and **one** capacitor
  - The bit is stored as a charge in the capacitor
  - Capacitor leaks charge over time
    - > Capacitors must be periodically recharged
    - > This is called **refresh**
    - > During refresh, DRAM can't be accessed
    - > Also after read, capacitor needs recharging again
  - Reading a DRAM cell is slower than reading SRAM
- **Pros:**
  - Higher density -> less silicon -> **much cheaper than SRAM**
- **Cons:**
  - Still volatile (even more volatile than SRAM)
  - **Slower access time**



# Spinning magnetic disks (HDD)

- Spinning platter coated with a ferromagnetic substance magnetized to represent bits
  - Has a mechanical arm with a head
  - Reads by placing arm in correct cylinder, and waiting for platter to rotate
- **Pros:**
  - **Nonvolatile** (magnetization persists without power)
  - Extremely cheap (1TB for \$50)
- **Cons:**
  - **Extremely slow** (it has a mechanical arm, enough said)



# Other technology

- Flash Memory
  - Works using a special MOSFET with “floating gate”
  - **Pros:** nonvolatile, much faster than HDD
  - **Cons:**
    - Slower than DRAM
    - More expensive than HDDs (1TB for \$250)
    - Writing is destructive and shortens lifespan
- Experimental technology
  - Ferroelectric RAM (FeRAM), Magnetoresistive RAM (MRAM), Phase-change memory (PRAM), carbon nanotubes ...
  - In varying states of development and maturity
  - Nonvolatile *and* close to DRAM speeds



# Do you notice a trend?

- The faster the memory the more expensive and lower density.
- The slower the memory the less expensive and higher density.
- Thus, memory is constructed as a hierarchy:
  - Fast and small memory at the upper levels
  - Slow and big memory at the lower levels
- And also data is stored hierarchically:
  - Frequently accessed data is stored in the fast, upper levels
  - Infrequently accessed data is stored in the slow, lower levels

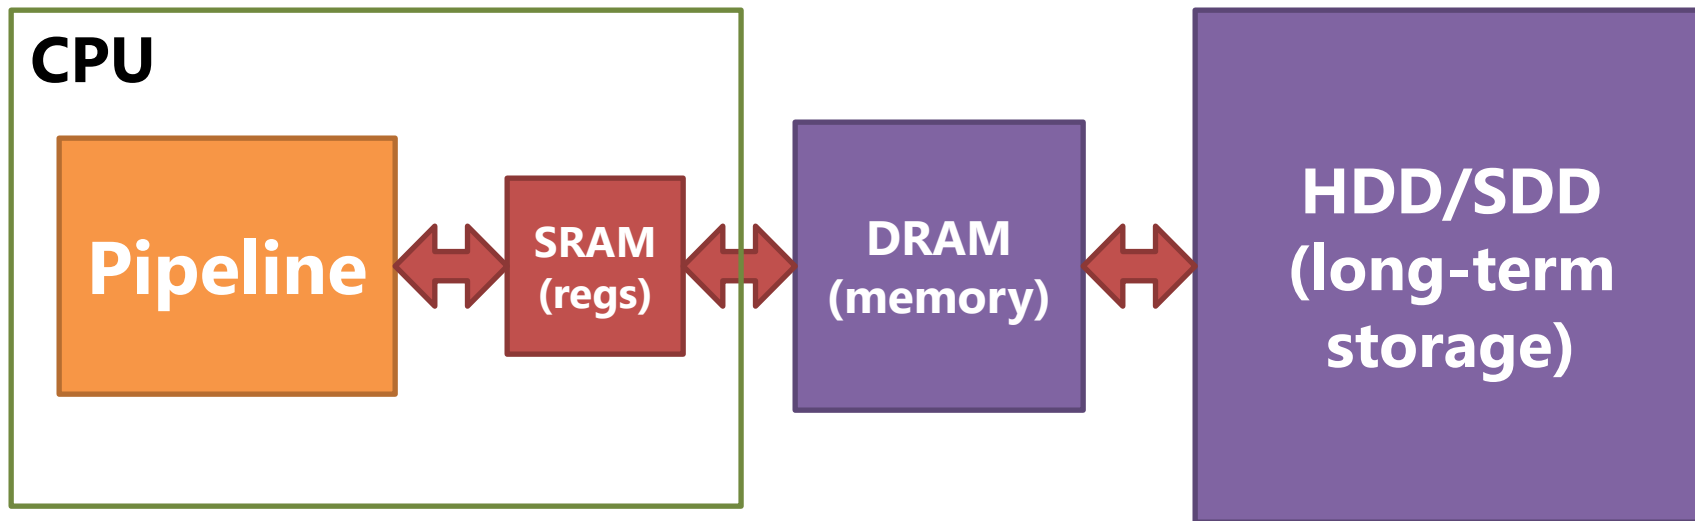
# The memory hierarchy

---



# From each according to their ability...

- Use the **fast memory closest to the CPU**, and use the **big memory furthest from it**.



- Speed tends to be inversely proportional to both price and size.

# Caching and Locality

- **Caching** is keeping a temporary copy of data for quick access.
- Each level in the hierarchy acts as a **cache** for the next lower level.
  - Registers are a temporary copy of data from memory.
  - Memory is a temporary copy of data from long-term storage.
  - Long-term storage is a temporary copy of... remote data?
- Caching exploits **locality**.
  - **Temporal locality** is the idea that you will access the same piece of data many times in succession.
  - **Spatial locality** is the idea that if you access a piece of data, you are likely to soon access another piece very close by
- By keeping data in a cache, we can greatly speed up data access by reusing the temporary copy.